

Introduction To Machine Learning: Project 2

SARIMZIA | 50245868

17 April 2018

Task 0

Preparations

Environment used: Rstudio Version 1.1.442 on Windows 10

TASK 1

8.3.1 Fitting classification Trees

```
library(tree)
```

```
## Warning: package 'tree' was built under R version 3.4.4
```

```
library(ISLR)
```

```
## Warning: package 'ISLR' was built under R version 3.4.4
```

```
attach(Carseats)
High <- with(Carseats, ifelse(Sales <= 8, "No", "Yes"))
Carseats <- data.frame(Carseats, High)
tree.carseats <- tree(High~.-Sales,Carseats)
summary(tree.carseats)
```

```
##
## Classification tree:
## tree(formula = High ~ . - Sales, data = Carseats)
## Variables actually used in tree construction:
## [1] "ShelveLoc"    "Price"        "Income"       "CompPrice"    "Population"
## [6] "Advertising"  "Age"          "US"
## Number of terminal nodes: 27
## Residual mean deviance: 0.4575 = 170.7 / 373
## Misclassification error rate: 0.09 = 36 / 400
```

```
plot(tree.carseats)
text(tree.carseats,pretty=0)
```



```
##          343) Age > 49.5 11    6.702 No ( 0.90909 0.09091 ) *
##          43) Price > 122.5 77   55.540 No ( 0.88312 0.11688 )
##          86) CompPrice < 147.5 58   17.400 No ( 0.96552 0.03448 ) *
##          87) CompPrice > 147.5 19   25.010 No ( 0.63158 0.36842 )
##          174) Price < 147 12    16.300 Yes ( 0.41667 0.58333 )
##          348) CompPrice < 152.5 7    5.742 Yes ( 0.14286 0.85714 ) *
##          349) CompPrice > 152.5 5    5.004 No ( 0.80000 0.20000 ) *
##          175) Price > 147 7    0.000 No ( 1.00000 0.00000 ) *
##        11) Advertising > 13.5 45   61.830 Yes ( 0.44444 0.55556 )
##          22) Age < 54.5 25   25.020 Yes ( 0.20000 0.80000 )
##          44) CompPrice < 130.5 14   18.250 Yes ( 0.35714 0.64286 )
##          88) Income < 100 9    12.370 No ( 0.55556 0.44444 ) *
##          89) Income > 100 5    0.000 Yes ( 0.00000 1.00000 ) *
##          45) CompPrice > 130.5 11    0.000 Yes ( 0.00000 1.00000 ) *
##          23) Age > 54.5 20   22.490 No ( 0.75000 0.25000 )
##          46) CompPrice < 122.5 10    0.000 No ( 1.00000 0.00000 ) *
##          47) CompPrice > 122.5 10   13.860 No ( 0.50000 0.50000 )
##          94) Price < 125 5    0.000 Yes ( 0.00000 1.00000 ) *
##          95) Price > 125 5    0.000 No ( 1.00000 0.00000 ) *
##        3) ShelveLoc: Good 85   90.330 Yes ( 0.22353 0.77647 )
##          6) Price < 135 68   49.260 Yes ( 0.11765 0.88235 )
##          12) US: No 17   22.070 Yes ( 0.35294 0.64706 )
##          24) Price < 109 8    0.000 Yes ( 0.00000 1.00000 ) *
##          25) Price > 109 9   11.460 No ( 0.66667 0.33333 ) *
##          13) US: Yes 51   16.880 Yes ( 0.03922 0.96078 ) *
##          7) Price > 135 17   22.070 No ( 0.64706 0.35294 )
##          14) Income < 46 6    0.000 No ( 1.00000 0.00000 ) *
##          15) Income > 46 11   15.160 Yes ( 0.45455 0.54545 ) *
```

```
set.seed(2)
train=sample(1:nrow(Carseats),200)
Carseats.test=Carseats[-train,]
High.test=High[-train]
tree.carseats=tree(High~.-Sales,Carseats,subset=train)
tree.pred=predict(tree.carseats,Carseats.test,type="class")
table(tree.pred,High.test)
```

```
##          High.test
## tree.pred No Yes
##          No  86  27
##          Yes 30  57
```

```
(86+57)/200
```

```
## [1] 0.715
```



```
set.seed(3)
cv.carseats=cv.tree(tree.carseats,FUN=prune.misclass)
names(cv.carseats)
```

```
## [1] "size" "dev" "k" "method"
```

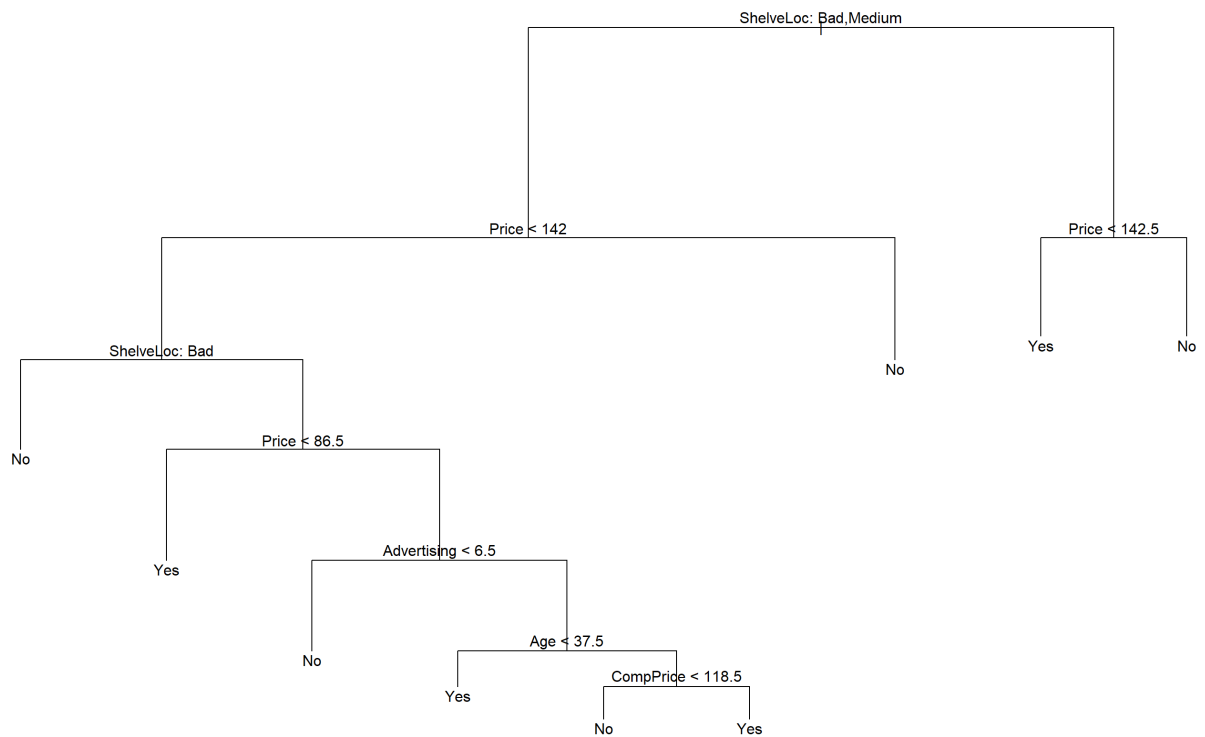
```
cv.carseats
```

```
## $size
## [1] 19 17 14 13 9 7 3 2 1
##
## $dev
## [1] 55 55 53 52 50 56 69 65 80
##
## $k
## [1] -Inf 0.0000000 0.6666667 1.0000000 1.7500000 2.0000000
## [7] 4.2500000 5.0000000 23.0000000
##
## $method
## [1] "misclass"
##
## attr(,"class")
## [1] "prune" "tree.sequence"
```

```
par(mfrow=c(1,2))
plot(cv.carseats$size,cv.carseats$dev,type="b")
plot(cv.carseats$k,cv.carseats$dev,type="b")
```



```
prune.carseats=prune.misclass(tree.carseats,best=9)
plot(prune.carseats)
text(prune.carseats,pretty=0)
```



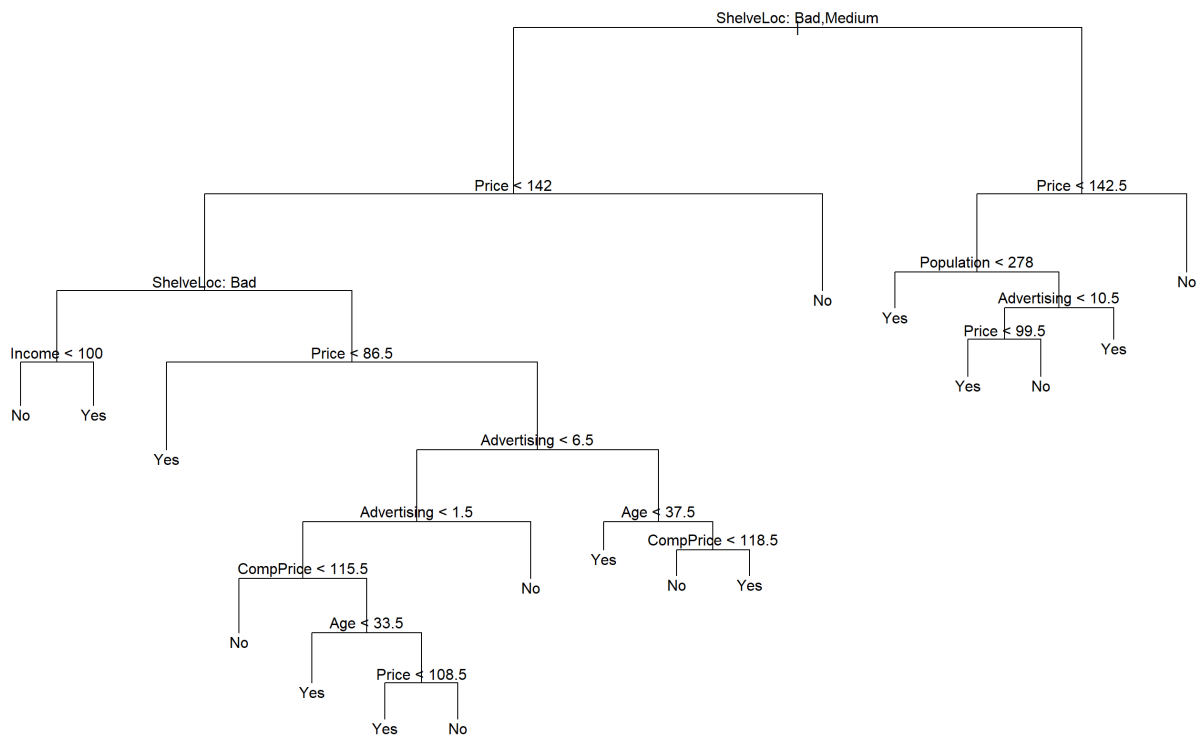
```
tree.pred=predict(prune.carseats,Carseats.test,type="class")
table(tree.pred,High.test)
```

```
##           High.test
## tree.pred No  Yes
##           No  94  24
##           Yes 22  60
```

```
(94+60)/200
```

```
## [1] 0.77
```

```
prune.carseats=prune.misclass(tree.carseats,best=15)
plot(prune.carseats)
text(prune.carseats,pretty=0)
```



```
tree.pred=predict(prune.carseats,Carseats.test,type="class")
table(tree.pred,High.test)
```

```
##           High.test
## tree.pred No  Yes
##           No   86  22
##           Yes  30  62
```

```
(86+62)/200
```

```
## [1] 0.74
```

TASK 2

8.3.2 Fitting Regression Trees


```
library(MASS)
```

```
## Warning: package 'MASS' was built under R version 3.4.4
```


```
set.seed=1
train=sample(1:nrow(Boston),nrow(Boston)/2)
tree.boston=tree(medv~.,Boston,subset=train)
summary(tree.boston)
```

```
##
## Regression tree:
## tree(formula = medv ~ ., data = Boston, subset = train)
## Variables actually used in tree construction:
## [1] "rm"      "lstat"    "dis"      "crim"     "ptratio"
## Number of terminal nodes: 9
## Residual mean deviance: 13.51 = 3298 / 244
## Distribution of residuals:
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -14.6900 -1.7790 -0.1793  0.0000  2.0210  17.5100
```


```
plot(tree.boston)
text(tree.boston, pretty=0)
```




```
cv.boston=cv.tree(tree.boston)
plot(cv.boston$size,cv.boston$dev,type='b')
```



```
prune.boston=prune.tree(tree.boston,best=5)
plot(prune.boston)
text(prune.boston,pretty=0)
```




```
yhat=predict(tree.boston,newdata=Boston[-train,])
boston.test=Boston[-train,"medv"]
plot(yhat,boston.test)
abline(0,1)
```



```
mean((yhat-boston.test)^2)
```

```
## [1] 24.48727
```

TASK 3

8.3.3 Bagging and Random Forests

```
library(randomForest)
```

```
## Warning: package 'randomForest' was built under R version 3.4.4
```


```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
set.seed(1)
bag.boston=randomForest(medv~.,data=Boston,subset=train,mtry=13,importance=TRUE)
bag.boston
```

```
##
## Call:
## randomForest(formula = medv ~ ., data = Boston, mtry = 13, importance = TRUE,
subset = train)
##
##           Type of random forest: regression
##           Number of trees: 500
## No. of variables tried at each split: 13
##
##           Mean of squared residuals: 11.78114
##           % Var explained: 85.82
```

```
yhat.bag=predict(bag.boston,newdata=Boston[-train,])
plot(yhat.bag, boston.test)
abline(0,1)
```



```
mean((yhat.bag-boston.test)^2)
```

```
## [1] 17.11319
```

```
bag.boston=randomForest(medv~.,data=Boston,subset=train,mtry=13,ntree=25)
yhat.bag=predict(bag.boston,newdata=Boston[-train,])
mean((yhat.bag-boston.test)^2)
```

```
## [1] 17.19022
```

```
set.seed(1)
rf.boston=randomForest(medv~.,data=Boston,subset=train,mtry=6,importance=TRUE)
yhat.rf=predict(rf.boston,newdata=Boston[-train,])
mean((yhat.rf-boston.test)^2)
```


```
## [1] 17.50849
```

```
importance(rf.boston)
```

```
##           %IncMSE  IncNodePurity
## crim      15.458985    1868.86996
## zn         4.381942      68.33245
## indus      9.014757     505.06808
## chas       2.047497      53.08569
## nox       13.085996    1130.86998
## rm        30.469371    6034.00836
## age       10.434793      733.86943
## dis       15.270284    1180.24510
## rad        4.830909      95.98200
## tax       12.299406     512.53066
## ptratio   10.686928      700.82904
## black      6.087007      285.82423
## lstat     31.377983     7286.11343
```

```
varImpPlot(rf.boston)
```

rf.boston



TASK 4

8.3.4 Boosting

```
library(gbm)
```

```
## Warning: package 'gbm' was built under R version 3.4.4
```

```
## Loading required package: survival
```


```
## Loading required package: lattice
```

```
## Loading required package: splines
```

```
## Loading required package: parallel
```



```
## Loaded gbm 2.1.3
```

```
set.seed(1)
boost.boston=gbm(medv~.,data=Boston[train,],distribution="gaussian",n.trees=5000,interaction.depth=4)
summary(boost.boston)
```



```
##           var      rel.inf
## lstat      lstat 50.16712978
## rm         rm   29.36863544
## dis        dis   5.27479136
## crim       crim   4.47470271
## nox        nox   2.84121915
## ptratio    ptratio 2.56803972
## tax        tax   1.98645724
## age        age   1.36264494
## indus      indus  0.97373021
## black      black  0.68537903
## rad        rad   0.22960165
## zn         zn    0.04429074
## chas       chas   0.02337802
```

```
par(mfrow=c(1,2))
plot(boost.boston,i="rm")
plot(boost.boston,i="lstat")
```



```
yhat.boost=predict(boost.boston,newdata=Boston[-train,],n.trees=5000)
mean((yhat.boost-boston.test)^2)
```

```
## [1] 17.47222
```

```
boost.boston=gbm(medv~.,data=Boston[train,],distribution ="gaussian",n.trees=5000,i
nteraction.depth = 4,shrinkage = 0.2,verbose=F)
yhat.boston=predict(boost.boston,newdata=Boston[-train,],n.trees=5000)
mean((yhat.boost-boston.test)^2)
```

```
## [1] 17.47222
```

TASK 5

Summary

It took me 2 days to finish this assignment. This time I did not encounter any difficulties as the tasks were pretty simple and the instructions were concise. I collaborated with Someshwar Rao for this assignment, only for the initial part, where I had to set up the environment and start with Rmarkdown.